

Chapter 1

Introduction to Machine Learning

Anaconda & Jupyter Installation


- Add Anaconda & Jupyter Installation Link
- Windows
- Mac
- Conda Install Seaborn

Python Packages and Libraries

- **Jupyter notebooks:** interactive coding and visualization of output
- **NumPy, SciPy, Pandas:** numerical computation
- **Matplotlib, Seaborn:** data visualization
- **Scikit-learn:** machine learning

Introduction to Jupyter Notebook


- Polyglot analysis environment—blends multiple languages



Source: <http://jupyter.org/>

Introduction to Jupyter Notebook

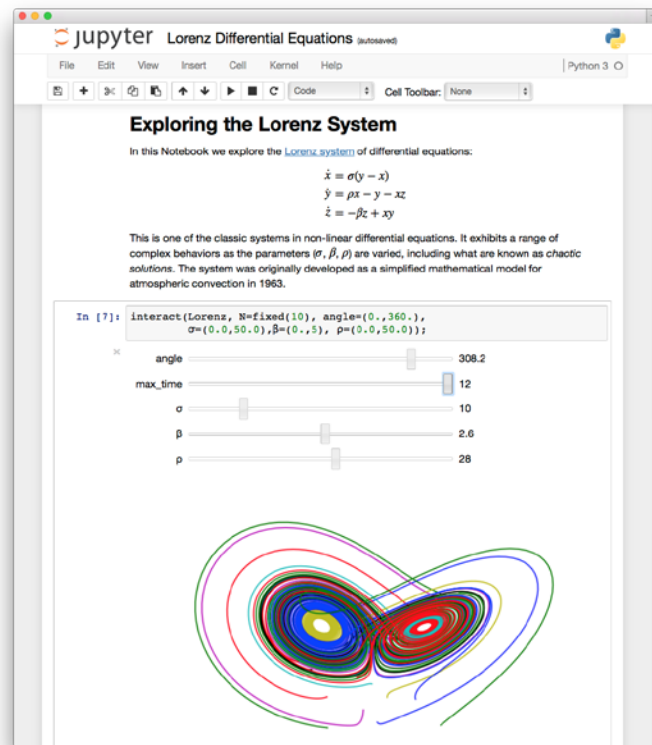
- Polyglot analysis environment—blends multiple languages
- Jupyter is an anagram of: Julia, Python, and R



Source: <http://jupyter.org/>

Introduction to Jupyter Notebook

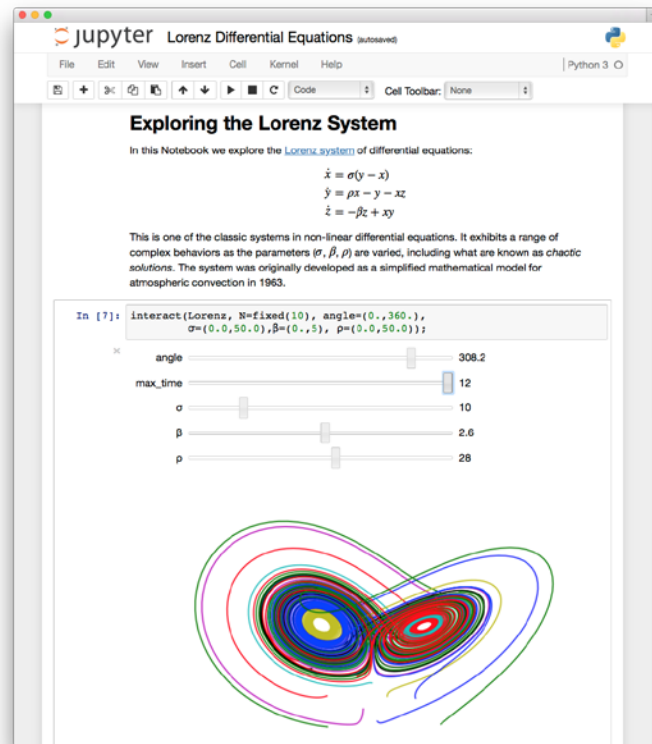
- Polyglot analysis environment—blends multiple languages
- Jupyter is an anagram of: Julia, Python, and R
- Supports multiple content types: code, narrative text, images, movies, etc.



Source: <http://jupyter.org/>

Introduction to Jupyter Notebook

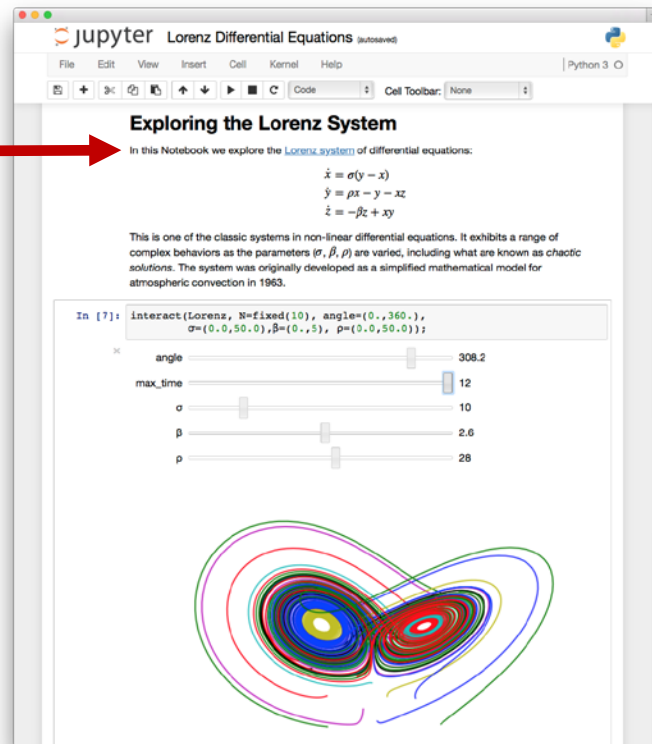
- HTML & Markdown
- LaTeX (equations)
- Code



Source: <http://jupyter.org/>

Introduction to Jupyter Notebook

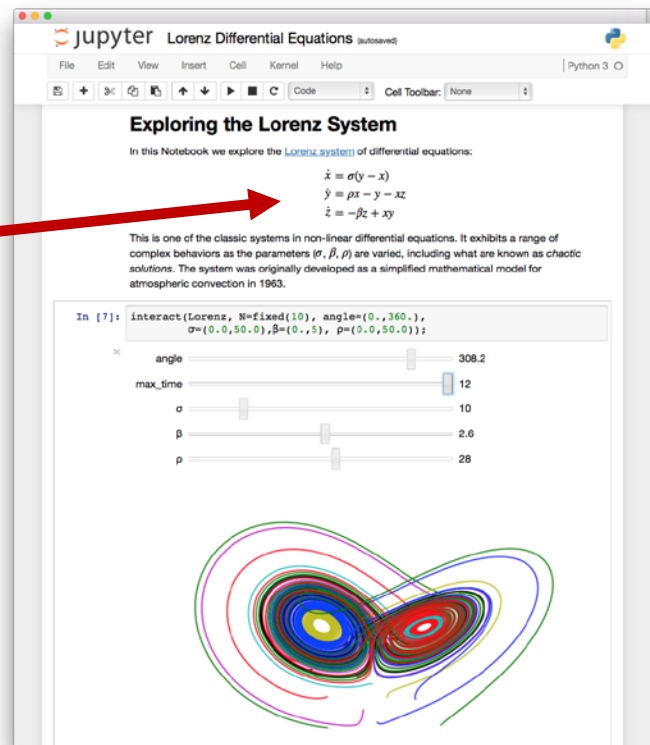
- HTML & Markdown
- LaTeX (equations)
- Code



Source: <http://jupyter.org/>

Introduction to Jupyter Notebook

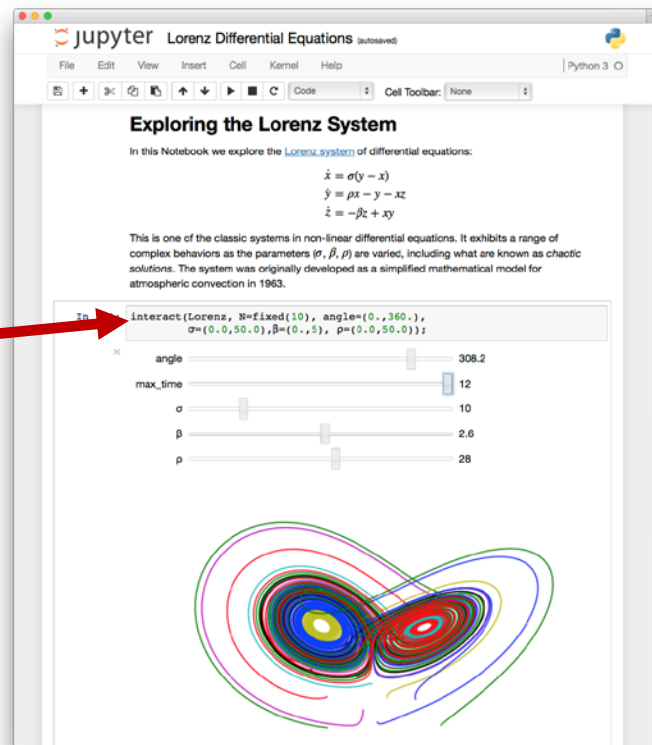
- HTML & Markdown
- LaTeX (equations)
- Code



Source: <http://jupyter.org/>

Introduction to Jupyter Notebook

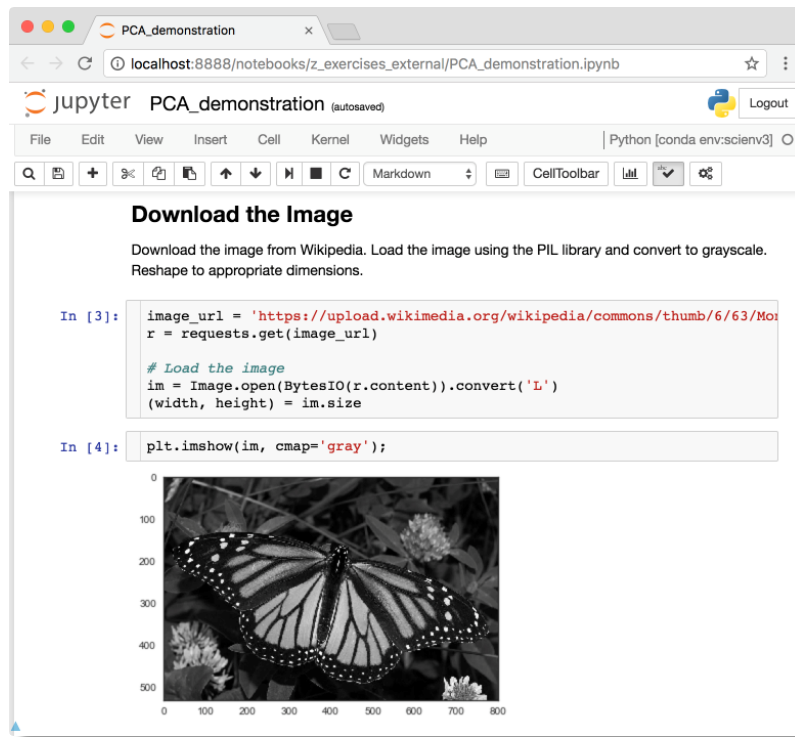
- HTML & Markdown
- LaTeX (equations)
- Code



Source: <http://jupyter.org/>

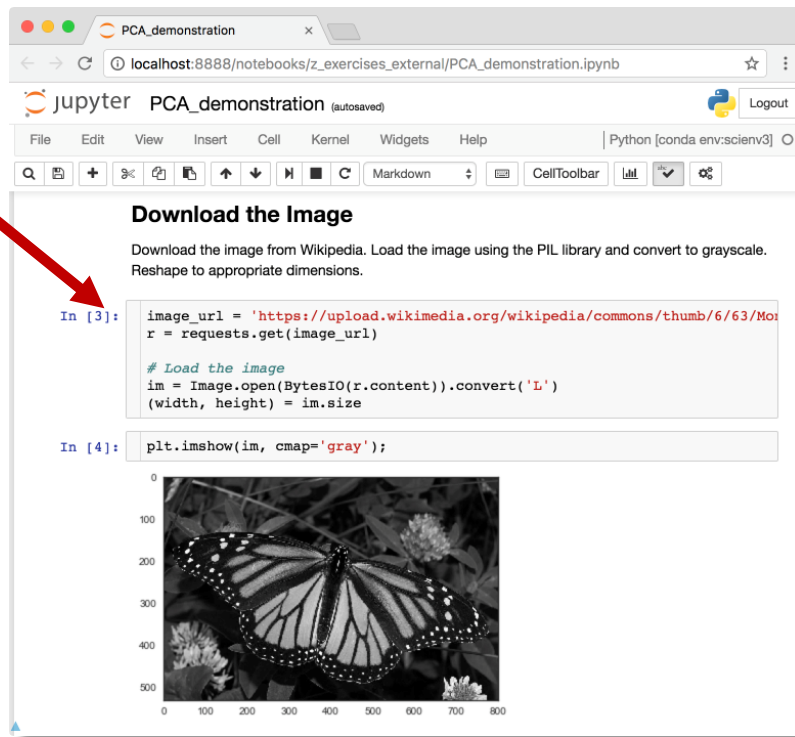
Introduction to Jupyter Notebook

- Code is divided into cells to control execution
- Enables interactive development
- Ideal for exploratory analysis and model building



Introduction to Jupyter Notebook

- Code is divided into cells to control execution
- Enables interactive development
- Ideal for exploratory analysis and model building



Jupyter Cell Magics

- `%matplotlib inline`: display plots inline in Jupyter notebook



The screenshot shows a Jupyter Notebook window titled "Jupyter Cell Magic Demo". The browser address bar indicates the URL is `localhost:8888/notebooks/z_exercises_external/Jupyter%20Cell%20Magic%20...`. The notebook interface includes a menu bar (File, Edit, View, Insert, Cell, Kernel, Widgets, Help) and a status bar showing "Python [conda env:scienv3]".

Three input cells are visible:

- In [1]:** Imports the necessary libraries:

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```
- In [2]:** Generates data and plots a line graph:

```
x = np.arange(10)
y = x**2
plt.plot(x, y);
```

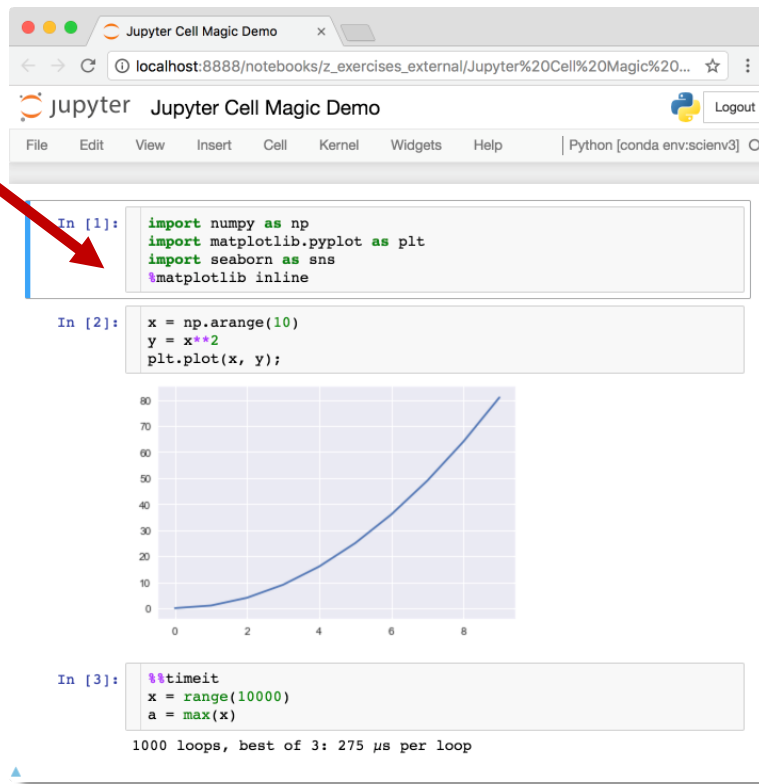
The plot shows a blue line representing the function $y = x^2$ for x values from 0 to 9. The x-axis is labeled from 0 to 8, and the y-axis is labeled from 0 to 80. The plot is displayed inline within the notebook cell.
- In [3]:** Uses the `%timeit` magic to measure the execution time of a loop:

```
%timeit
x = range(10000)
a = max(x)
```

The output of this cell is: "1000 loops, best of 3: 275 µs per loop".

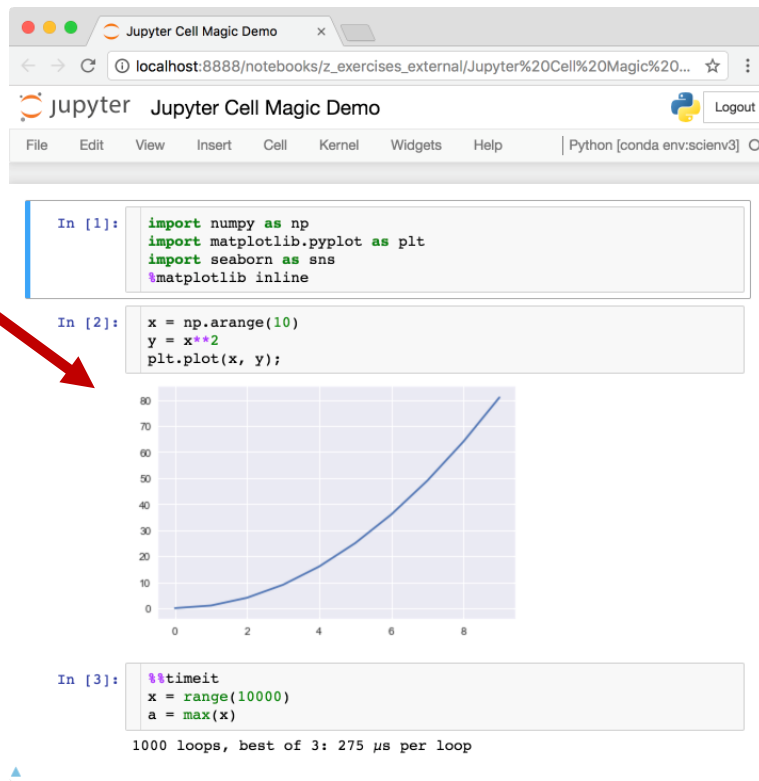
Jupyter Cell Magics

- `%matplotlib inline`: display plots inline in Jupyter notebook



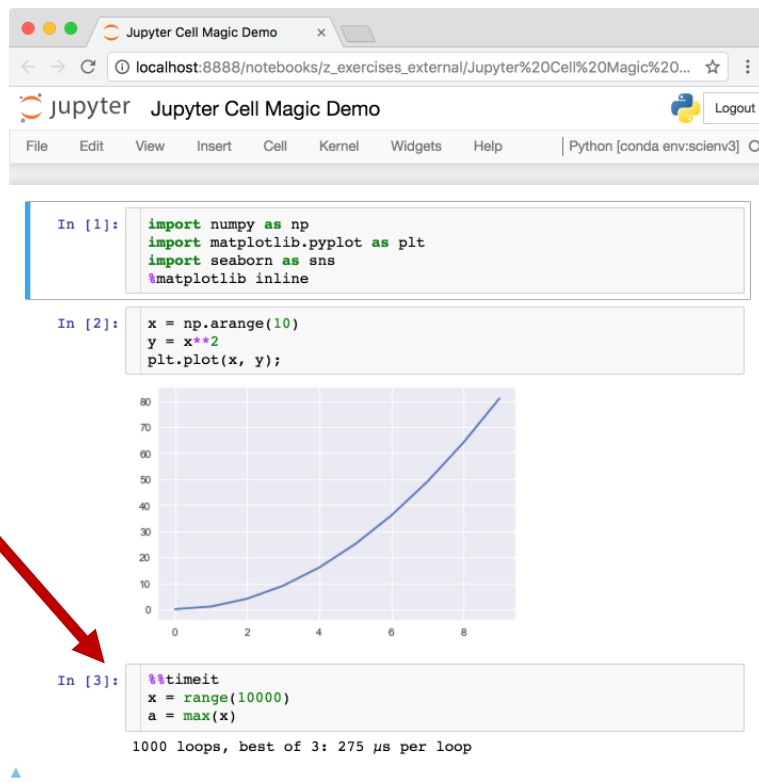
Jupyter Cell Magics

- `%matplotlib inline`: display plots inline in Jupyter notebook



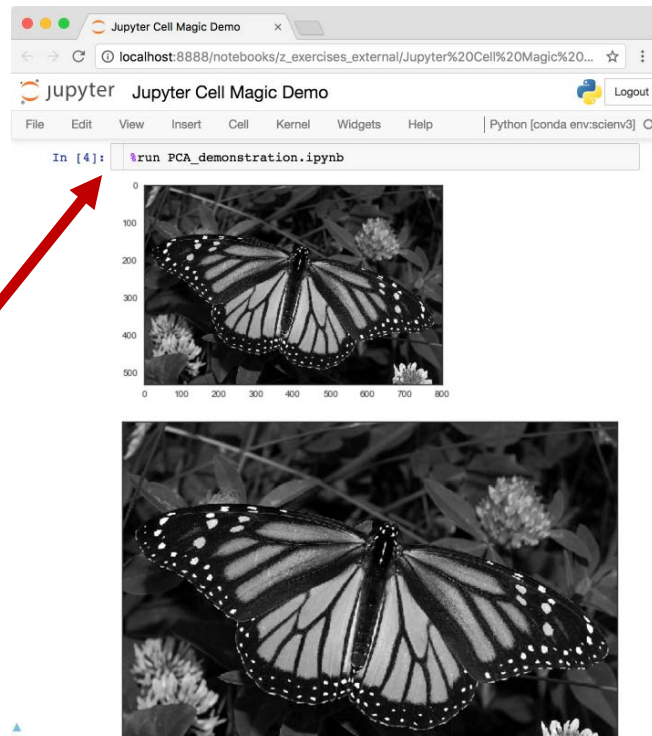
Jupyter Cell Magics

- `%matplotlib inline`: display plots inline in Jupyter notebook
- `%%timeit`: time how long a cell takes to execute



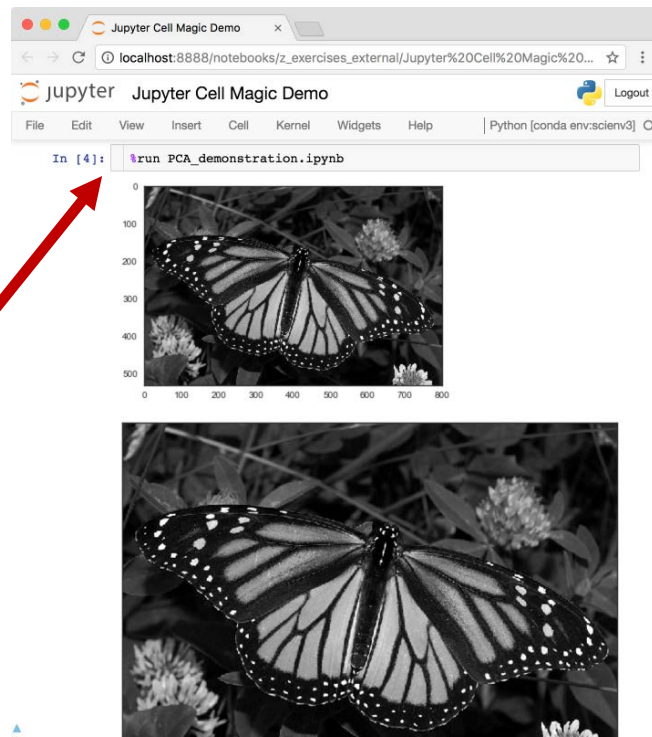
Jupyter Cell Magics

- `%matplotlib inline`: display plots inline in Jupyter notebook
- `%%timeit`: time how long a cell takes to execute
- `%run filename.ipynb`: execute code from another notebook or python file



Jupyter Cell Magics

- `%matplotlib inline`: display plots inline in Jupyter notebook
- `%%timeit`: time how long a cell takes to execute
- `%run filename.ipynb`: execute code from another notebook or python file
- `%load filename.py`: copy contents of the file and paste into the cell



Jupyter Keyboard Shortcuts

Keyboard shortcuts

The Jupyter Notebook has two different keyboard input modes. **Edit mode** allows you to type code/text into a cell and is indicated by a green cell border. **Command mode** binds the keyboard to notebook level actions and is indicated by a grey cell border with a blue left margin.

Command Mode (press `Esc` to enable)

`F`: find and replace

`Ctrl-Shift-P`: open the command palette

`Enter`: enter edit mode

`Shift-Enter`: run cell, select below

`Ctrl-Enter`: run selected cells

`Alt-Enter`: run cell, insert below

`Shift-J`: extend selected cells below

`A`: insert cell above

`B`: insert cell below

`X`: cut selected cells

`C`: copy selected cells

`Shift-V`: paste cells above

Keyboard shortcuts can be viewed from Help → Keyboard Shortcuts

Making Jupyter Notebooks Reusable

To extract Python code from a Jupyter notebook:

Convert from Command Line

```
>>> jupyter nbconvert --to python  
notebook.ipynb
```

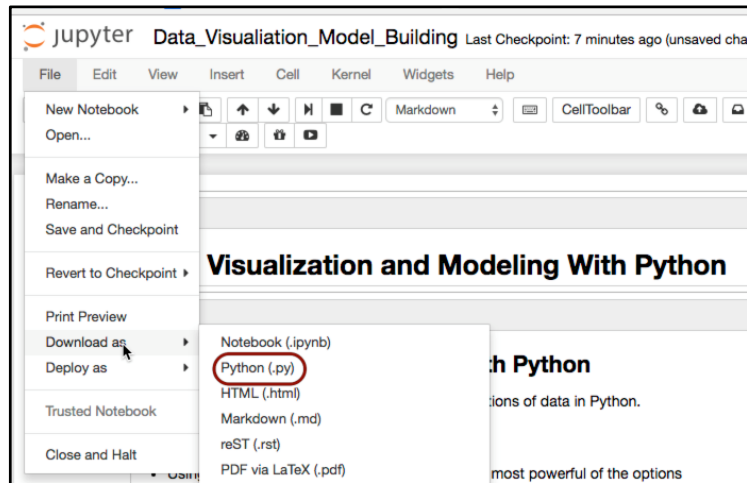
Making Jupyter Notebooks Reusable

To extract Python code from a Jupyter notebook:

Convert from Command Line

```
>>> jupyter nbconvert --to python  
notebook.ipynb
```

Export from Notebook




Introduction to Pandas

- Library for computation with tabular data
- Mixed types of data allowed in a single table
- Columns and rows of data can be named
- Advanced data aggregation and statistical functions

pandas


$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



Source: <http://pandas.pydata.org/>

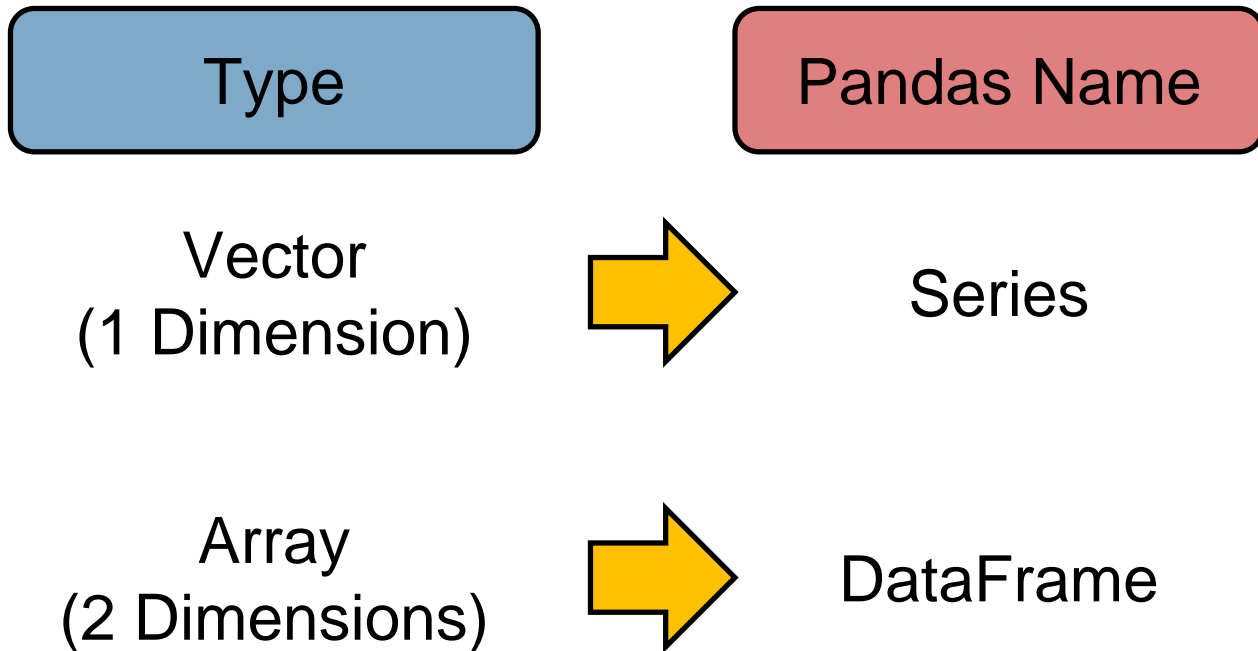
Introduction to Pandas

Basic data structures



Introduction to Pandas

Basic data structures



Pandas Series Creation and Indexing

Use data from step tracking application to create a Pandas Series

Code

```
import pandas as pd

step_data = [3620, 7891, 9761,
             3907, 4338, 5373]

step_counts = pd.Series(step_data,
                        name='steps')

print(step_counts)
```

Output

Pandas Series Creation and Indexing

Use data from step tracking application to create a Pandas Series

Code

```
import pandas as pd

step_data = [3620, 7891, 9761,
             3907, 4338, 5373]

step_counts = pd.Series(step_data,
                        name='steps')

print(step_counts)
```

Output

```
>>> 0 3620
      1 7891
      2 9761
      3 3907
      4 4338
      5 5373
      Name: steps, dtype: int64
```

Pandas Series Creation and Indexing

Add a date range to the Series

Code

```
step_counts.index = pd.date_range('20150329',  
                                   periods=6)  
  
print(step_counts)
```

Output

Pandas Series Creation and Indexing

Add a date range to the Series

Code

```
step_counts.index = pd.date_range('20150329',  
                                   periods=6)  
  
print(step_counts)
```

Output

```
>>> 2015-03-29 3620  
      2015-03-30 7891  
      2015-03-31 9761  
      2015-04-01 3907  
      2015-04-02 4338  
      2015-04-03 5373  
      Freq: D, Name: steps,  
      dtype: int64
```

Pandas Series Creation and Indexing

Select data by the index values

Code

```
# Just like a dictionary  
print(step_counts['2015-04-01'])
```

Output

Pandas Series Creation and Indexing

Select data by the index values

Code

```
# Just like a dictionary  
print(step_counts['2015-04-01'])
```

Output

```
>>> 3907
```

Pandas Series Creation and Indexing

Select data by the index values

Code

```
# Just like a dictionary  
print(step_counts['2015-04-01'])  
  
# Or by index position--like an array  
print(step_counts[3])
```

Output

```
>>> 3907
```

Pandas Series Creation and Indexing

Select data by the index values

Code

```
# Just like a dictionary  
print(step_counts['2015-04-01'])  
  
# Or by index position--like an array  
print(step_counts[3])
```

Output

```
>>> 3907
```

```
>>> 3907
```


Pandas Series Creation and Indexing

Select data by the index values

Code

```
# Just like a dictionary
print(step_counts['2015-04-01'])

# Or by index position--like an array
print(step_counts[3])

# Select all of April
print(step_counts['2015-04'])
```

Output

```
>>> 3907
```

```
>>> 3907
```

Pandas Series Creation and Indexing

Select data by the index values

Code

```
# Just like a dictionary
print(step_counts['2015-04-01'])

# Or by index position--like an array
print(step_counts[3])

# Select all of April
print(step_counts['2015-04'])
```

Output

```
>>> 3907
```

```
>>> 3907
```

```
>>> 2015-04-01 3907
      2015-04-02 4338
      2015-04-03 5373
      Freq: D, Name: steps,
      dtype: int64
```

Pandas Data Types and Imputation

Data types can be viewed and converted

Code

```
# View the data type  
print(step_counts.dtypes)
```

Output

Pandas Data Types and Imputation

Data types can be viewed and converted

Code

```
# View the data type  
print(step_counts.dtypes)
```

Output

```
>>> int64
```

Pandas Data Types and Imputation

Data types can be viewed and converted

Code

```
# View the data type
print(step_counts.dtypes)

# Convert to a float
step_counts = step_counts.astype(np.float)

# View the data type
print(step_counts.dtypes)
```

Output

```
>>> int64
```

Pandas Data Types and Imputation

Data types can be viewed and converted

Code

```
# View the data type
print(step_counts.dtypes)

# Convert to a float
step_counts = step_counts.astype(np.float)

# View the data type
print(step_counts.dtypes)
```

Output

```
>>> int64
```

```
>>> float64
```

Pandas Data Types and Imputation

Invalid data points can be easily filled with values

Code

```
# Create invalid data
step_counts[1:3] = np.NaN

# Now fill it in with zeros
step_counts = step_counts.fillna(0.)
# equivalently,
# step_counts.fillna(0., inplace=True)

print(step_counts[1:3])
```

Output

Pandas Data Types and Imputation

Invalid data points can be easily filled with values

Code

```
# Create invalid data
step_counts[1:3] = np.NaN

# Now fill it in with zeros
step_counts = step_counts.fillna(0.)
# equivalently,
# step_counts.fillna(0., inplace=True)

print(step_counts[1:3])
```

Output

```
>>> 2015-03-30  0.0
      2015-03-31  0.0
      Freq: D, Name: steps,
      dtype: float64
```


Pandas DataFrame Creation and Methods

DataFrames can be created from lists, dictionaries, and Pandas Series

Code

```
# Cycling distance
cycling_data = [10.7, 0, None, 2.4, 15.3,
                10.9, 0, None]

# Create a tuple of data
joined_data = list(zip(step_data,
                       cycling_data))

# The dataframe
activity_df = pd.DataFrame(joined_data)

print(activity_df)
```

Output

Pandas DataFrame Creation and Methods

DataFrames can be created from lists, dictionaries, and Pandas Series

Code

```
# Cycling distance
cycling_data = [10.7, 0, None, 2.4, 15.3,
                10.9, 0, None]

# Create a tuple of data
joined_data = list(zip(step_data,
                       cycling_data))

# The dataframe
activity_df = pd.DataFrame(joined_data)

print(activity_df)
```

Output

>>>

	0	1
0	3620	10.7
1	7891	0.0
2	9761	NaN
3	3907	2.4
4	4338	15.3
5	5373	10.9

Pandas DataFrame Creation and Methods

Labeled columns and an index can be added

Code

```
# Add column names to dataframe
activity_df = pd.DataFrame(
    joined_data,
    index=pd.date_range('20150329', periods=6),
    columns=['Walking', 'Cycling'])

print(activity_df)
```

Output

Pandas DataFrame Creation and Methods

Labeled columns and an index can be added

Code

```
# Add column names to dataframe
activity_df = pd.DataFrame(joined_data,
                           index=pd.date_range('20150329',
                                                periods=6),
                           columns=['Walking', 'Cycling'])

print(activity_df)
```

Output

>>>

	Walking	Cycling
2015-03-29	3620	10.7
2015-03-30	7891	0.0
2015-03-31	9761	NaN
2015-04-01	3907	2.4
2015-04-02	4338	15.3
2015-04-03	5373	10.9

Indexing DataFrame Rows

DataFrame rows can be indexed by row using the 'loc' and 'iloc' methods

Code

```
# Select row of data by index name  
print(activity_df.loc['2015-04-01'])
```

Output

Indexing DataFrame Rows

DataFrame rows can be indexed by row using the 'loc' and 'iloc' methods

Code

```
# Select row of data by index name  
print(activity_df.loc['2015-04-01'])
```

Output

```
>>> Walking 3907.0  
      Cycling 2.4  
      Name: 2015-04-01,  
      dtype: float64
```

Indexing DataFrame Rows

DataFrame rows can be indexed by row using the 'loc' and 'iloc' methods

Code

```
# Select row of data by integer position  
print(activity_df.iloc[-3])
```

Output

Indexing DataFrame Rows

DataFrame rows can be indexed by row using the 'loc' and 'iloc' methods

Code

```
# Select row of data by integer position  
print(activity_df.iloc[-3])
```

Output

```
>>> Walking 3907.0  
      Cycling 2.4  
      Name: 2015-04-01,  
      dtype: float64
```


Indexing DataFrame Columns

DataFrame columns can be indexed by name

Code

```
# Name of column  
print(activity_df['Walking'])
```

Output

Indexing DataFrame Columns

DataFrame columns can be indexed by name

Code

```
# Name of column  
print(activity_df['Walking'])
```

Output

```
>>> 2015-03-29 3620  
      2015-03-30 7891  
      2015-03-31 9761  
      2015-04-01 3907  
      2015-04-02 4338  
      2015-04-03 5373  
Freq: D, Name: Walking,  
dtype: int64
```

Indexing DataFrame Columns

DataFrame columns can also be indexed as properties

Code

```
# Object-oriented approach  
print(activity_df.Walking)
```

Output

Indexing DataFrame Columns

DataFrame columns can also be indexed as properties

Code

```
# Object-oriented approach  
print(activity_df.Walking)
```

Output

```
>>> 2015-03-29 3620  
      2015-03-30 7891  
      2015-03-31 9761  
      2015-04-01 3907  
      2015-04-02 4338  
      2015-04-03 5373  
      Freq: D, Name: Walking,  
      dtype: int64
```

Indexing DataFrame Columns

DataFrame columns can be indexed by integer

Code

```
# First column  
print(activity_df.iloc[:,0])
```

Output

Indexing DataFrame Columns

DataFrame columns can be indexed by integer

Code

```
# First column  
print(activity_df.iloc[:,0])
```

Output

```
>>> 2015-03-29 3620  
      2015-03-30 7891  
      2015-03-31 9761  
      2015-04-01 3907  
      2015-04-02 4338  
      2015-04-03 5373  
      Freq: D, Name: Walking,  
      dtype: int64
```

Reading Data with Pandas

CSV and other common filetypes can be read with a single command

Code

```
# The location of the data file
filepath = 'data/Iris_Data/Iris_Data.csv'

# Import the data
data = pd.read_csv(filepath)

# Print a few rows
print(data.iloc[:5])
```

Output

Reading Data with Pandas

CSV and other common filetypes can be read with a single command

Code

```
# The location of the data file
filepath = 'data/Iris_Data/Iris_Data.csv'

# Import the data
data = pd.read_csv(filepath)

# Print a few rows
print(data.iloc[:5])
```

Output

>>>

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

Assigning New Data to a DataFrame

Data can be (re-)assigned to a DataFrame column

Code

```
# Create a new column that is a product  
# of both measurements  
data['sepal_area'] = data.sepal_length *  
                    data.sepal_width  
  
# Print a few rows and columns  
print(data.iloc[:5, -3:])
```

Output

Assigning New Data to a DataFrame

Data can be (re-)assigned to a DataFrame column

Code

```
# Create a new column that is a product  
# of both measurements  
data['sepal_area'] = data.sepal_length *  
                    data.sepal_width  
  
# Print a few rows and columns  
print(data.iloc[:5, -3:])
```

Output

>>>

	petal_width	species	sepal_area
0	0.2	Iris-setosa	17.85
1	0.2	Iris-setosa	14.70
2	0.2	Iris-setosa	15.04
3	0.2	Iris-setosa	14.26
4	0.2	Iris-setosa	18.00

Applying a Function to a DataFrame Column

Functions can be applied to columns or rows of a DataFrame or Series

Code

```
# The lambda function applies what  
# follows it to each row of data  
data['abbrev'] = (data  
                  .species  
                  .apply(lambda x:  
                        x.replace('Iris-', '')))  
  
# Note that there are other ways to  
# accomplish the above  
  
print(data.iloc[:5, -3:])
```

Output

Applying a Function to a DataFrame Column

Functions can be applied to columns or rows of a DataFrame or Series

Code

```
# The lambda function applies what  
# follows it to each row of data  
data['abbrev'] = (data  
                  .species  
                  .apply(lambda x:  
                        x.replace('Iris-', '')))  
  
# Note that there are other ways to  
# accomplish the above  
  
print(data.iloc[:5, -3:])
```

Output

```
>>>
```

	petal_width	species	abbrev
0	0.2	Iris-setosa	setosa
1	0.2	Iris-setosa	setosa
2	0.2	Iris-setosa	setosa
3	0.2	Iris-setosa	setosa
4	0.2	Iris-setosa	setosa

Concatenating Two DataFrames

Two DataFrames can be concatenated along either dimension

Code

```
# Concatenate the first two and  
# last two rows  
small_data = pd.concat([data.iloc[:2],  
                        data.iloc[-2:]])  
  
print(small_data.iloc[:,-3:])  
  
# See the 'join' method for  
# SQL style joining of dataframes
```

Output

Concatenating Two DataFrames

Two DataFrames can be concatenated along either dimension

Code

```
# Concatenate the first two and  
# last two rows  
small_data = pd.concat([data.iloc[:2],  
                        data.iloc[-2:]])  
  
print(small_data.iloc[:,-3:])  
  
# See the 'join' method for  
# SQL style joining of dataframes
```

Output

>>>

	petal_length	petal_width	species
0	1.4	0.2	Iris-setosa
1	1.4	0.2	Iris-setosa
148	5.4	2.3	Iris-virginica
149	5.1	1.8	Iris-virginica

Aggregated Statistics with GroupBy

Using the groupby method calculated aggregated DataFrame statistics

Code

```
# Use the size method with a  
# DataFrame to get count  
# For a Series, use the .value_counts  
# method  
group_sizes = (data  
                .groupby('species')  
                .size())  
  
print(group_sizes)
```

Output

Aggregated Statistics with GroupBy

Using the groupby method calculated aggregated DataFrame statistics

Code

```
# Use the size method with a  
# DataFrame to get count  
# For a Series, use the .value_counts  
# method  
group_sizes = (data  
                .groupby('species')  
                .size())  
  
print(group_sizes)
```

Output

```
>>> species  
      Iris-setosa      50  
      Iris-versicolor  50  
      Iris-virginica   50  
      dtype: int64
```


Performing Statistical Calculations

Pandas contains a variety of statistical methods—mean, median, and mode

Code

```
# Mean calculated on a DataFrame  
print(data.mean())
```

Output

Performing Statistical Calculations

Pandas contains a variety of statistical methods—mean, median, and mode

Code

```
# Mean calculated on a DataFrame  
print(data.mean())
```

Output

```
>>> sepal_length 5.843333  
      sepal_width 3.054000  
      petal_length 3.758667  
      petal_width 1.198667  
      dtype: float64
```

Performing Statistical Calculations

Pandas contains a variety of statistical methods—mean, median, and mode

Code

```
# Mean calculated on a DataFrame
```

```
print(data.mean())
```

```
# Median calculated on a Series
```

```
print(data.petal_length.median())
```

Output

```
>>> sepal_length 5.843333  
      sepal_width 3.054000  
      petal_length 3.758667  
      petal_width 1.198667  
      dtype: float64
```

```
>>> 4.35
```

Performing Statistical Calculations

Pandas contains a variety of statistical methods—mean, median, and mode

Code

```
# Mean calculated on a DataFrame
```

```
print(data.mean())
```

```
# Median calculated on a Series
```

```
print(data.petal_length.median())
```

```
# Mode calculated on a Series
```

```
print(data.petal_length.mode())
```

Output

```
>>> sepal_length 5.843333  
      sepal_width 3.054000  
      petal_length 3.758667  
      petal_width 1.198667  
      dtype: float64
```

```
>>> 4.35
```

```
>>> 0 1.5  
      dtype: float64
```

Performing Statistical Calculations

Standard deviation, variance, SEM and quantiles can also be calculated

Code

```
# Standard dev, variance, and SEM  
print(data.petal_length.std(),  
      data.petal_length.var(),  
      data.petal_length.sem())
```

Output

Performing Statistical Calculations

Standard deviation, variance, SEM and quantiles can also be calculated

Code

```
# Standard dev, variance, and SEM  
print(data.petal_length.std(),  
      data.petal_length.var(),  
      data.petal_length.sem())
```

Output

```
>>> 1.76442041995  
      3.11317941834  
      0.144064324021
```

Performing Statistical Calculations

Standard deviation, variance, SEM and quantiles can also be calculated

Code

```
# Standard dev, variance, and SEM
```

```
print(data.petal_length.std(),  
      data.petal_length.var(),  
      data.petal_length.sem())
```

```
# As well as quantiles
```

```
print(data.quantile(0))
```

Output

```
>>> 1.76442041995  
      3.11317941834  
      0.144064324021
```

```
>>> sepal_length 4.3  
      sepal_width 2.0  
      petal_length 1.0  
      petal_width 0.1  
      Name: 0, dtype: float64
```

Performing Statistical Calculations

Multiple calculations can be presented in a DataFrame

Code

```
print(data.describe())
```

Output

Performing Statistical Calculations

Multiple calculations can be presented in a DataFrame

Code

```
print(data.describe())
```

Output

```
>>>
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

Sampling from DataFrames

DataFrames can be randomly sampled from

Code

```
#Sample 5 rows without replacement
sample = (data
          .sample(n=5,
                  replace=False,
                  random_state=42))

print(sample.iloc[:, -3:])
```

Output

Sampling from DataFrames

DataFrames can be randomly sampled from

Code

```
#Sample 5 rows without replacement
sample = (data
          .sample(n=5,
                  replace=False,
                  random_state=42))

print(sample.iloc[:, -3:])
```

Output

>>>

	petal_length	petal_width	species
73	4.7	1.2	Iris-versicolor
18	1.7	0.3	Iris-setosa
118	6.9	2.3	Iris-virginica
78	4.5	1.5	Iris-versicolor
76	4.8	1.4	Iris-versicolor

Sampling from DataFrames

DataFrames can be randomly sampled from

Code

```
#Sample 5 rows without replacement
sample = (data
          .sample(n=5,
                  replace=False,
                  random_state=42))

print(sample.iloc[:, -3:])
```

Output

>>>

	petal_length	petal_width	species
73	4.7	1.2	Iris-versicolor
18	1.7	0.3	Iris-setosa
118	6.9	2.3	Iris-virginica
78	4.5	1.5	Iris-versicolor
76	4.8	1.4	Iris-versicolor

SciPy and NumPy also contain a variety of statistical functions.

Visualization Libraries

Visualizations can be created in multiple ways:

- Matplotlib
- Pandas (via Matplotlib)
- Seaborn
 - Statistically-focused plotting methods
 - Global preferences incorporated by Matplotlib

Basic Scatter Plots with Matplotlib

Scatter plots can be created from Pandas Series

Code

```
Import matplotlib.pyplot as plt
```

```
plt.plot(data.sepal_length,  
         data.sepal_width,  
         ls = ' ', marker='o')
```

Output

Basic Scatter Plots with Matplotlib


Scatter plots can be created from Pandas Series

Code

```
Import matplotlib.pyplot as plt
```

```
plt.plot(data.sepal_length,  
         data.sepal_width,  
         ls='', marker='o')
```

Output



Basic Scatter Plots with Matplotlib

Multiple layers of data can also be added

Code

```
plt.plot(data.sepal_length,  
         data.sepal_width,  
         ls='', marker='o',  
         label='sepal')  
  
plt.plot(data.petal_length,  
         data.petal_width,  
         ls='', marker='o',  
         label='petal')
```

Output


Basic Scatter Plots with Matplotlib

Multiple layers of data can also be added

Code

```
plt.plot(data.sepal_length,  
         data.sepal_width,  
         ls='', marker='o',  
         label='sepal')  
  
plt.plot(data.petal_length,  
         data.petal_width,  
         ls='', marker='o',  
         label='petal')
```

Output



Histograms with Matplotlib

Histograms can be created from Pandas Series

Code

```
plt.hist(data.sepal_length, bins=25)
```

Output


Histograms with Matplotlib

Histograms can be created from Pandas Series

Code

```
plt.hist(data.sepal_length, bins=25)
```

Output



Customizing Matplotlib Plots

Every feature of Matplotlib plots can be customized

Code

```
fig, ax = plt.subplots()

ax.barh(np.arange(10),
        data.sepal_width.iloc[:10])

# Set position of ticks and tick labels
ax.set_yticks(np.arange(0.4,10.4,1.0))
ax.set_yticklabels(np.arange(1,11))
ax.set(xlabel='xlabel', ylabel='ylabel',
       title='Title')
```

Output

Customizing Matplotlib Plots

Every feature of Matplotlib plots can be customized


Code

```
fig, ax = plt.subplots()

ax.barh(np.arange(10),
        data.sepal_width.iloc[:10])

# Set position of ticks and tick labels
ax.set_yticks(np.arange(0.4,10.4,1.0))
ax.set_yticklabels(np.arange(1,11))
ax.set(xlabel='xlabel', ylabel='ylabel',
       title='Title')
```

Output



Incorporating Statistical Calculations

Statistical calculations can be included with Pandas methods

Code

```
(data
 .groupby('species')
 .mean()
 .plot(color=['red','blue',
             'black','green'],
        fontsize=10.0, figsize=(4,4)))
```

Output

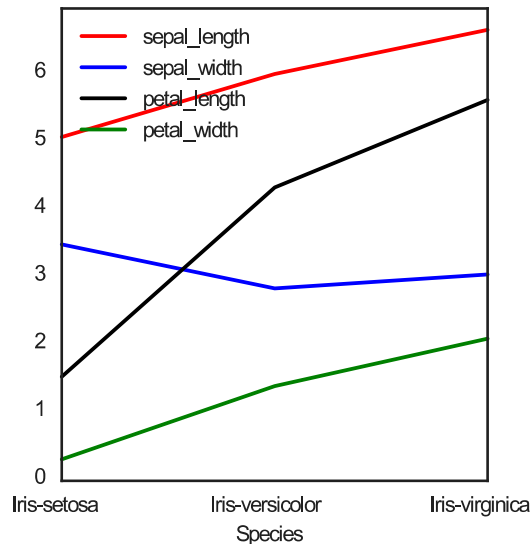
Incorporating Statistical Calculations

Statistical calculations can be included with Pandas methods

Code

```
(data
.groupby('species')
.mean()
.plot(color=['red','blue',
            'black','green'],
      fontsize=10.0, figsize=(4,4)))
```

Output



Statistical Plotting with Seaborn

Joint distribution and scatter plots can be created

Code

```
import seaborn as sns

sns.jointplot(x='sepal_length',
              y='sepal_width',
              data=data, size=4)
```

Output

Statistical Plotting with Seaborn


Joint distribution and scatter plots can be created

Code

```
import seaborn as sns

sns.jointplot(x='sepal_length',
              y='sepal_width',
              data=data, size=4)
```

Output



Statistical Plotting with Seaborn

Correlation plots of all variable pairs can also be made with Seaborn

Code

```
sns.pairplot(data, hue='species', size=3)
```

Output

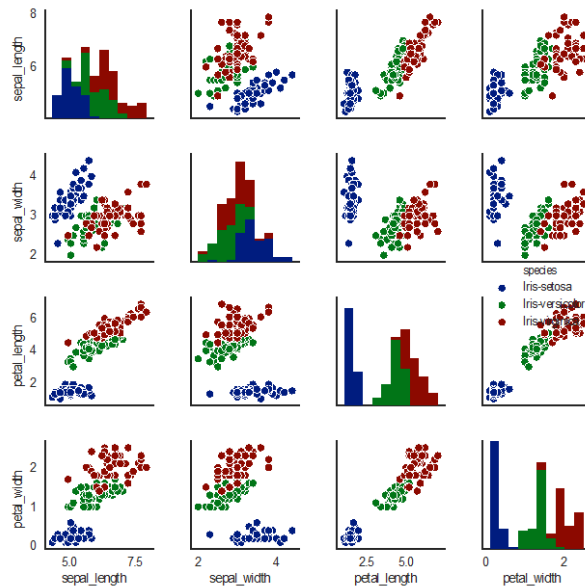
Statistical Plotting with Seaborn

Correlation plots of all variable pairs can also be made with Seaborn

Code

```
sns.pairplot(data, hue='species', size=3)
```

Output



Legal Notices and Disclaimers

This presentation is for informational purposes only. INTEL MAKES NO WARRANTIES, EXPRESS OR IMPLIED, IN THIS SUMMARY.

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. Check with your system manufacturer or retailer or learn more at [intel.com](https://www.intel.com).

This sample source code is released under the [Intel Sample Source Code License Agreement](#).

Intel and the Intel logo are trademarks of Intel Corporation in the U.S. and/or other countries.

*Other names and brands may be claimed as the property of others.

Copyright © 2017, Intel Corporation. All rights reserved.