


Chapter 7

Support Vector Machines Kernels


Support Vector Machines

Relationship to Logistic Regression




$$y_{\beta}(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x + \varepsilon)}}$$

Support Vector Machines (SVM)




Support Vector Machines (SVM)




Three
misclassifications

Support Vector Machines (SVM)




Two misclassifications

Support Vector Machines (SVM)




No misclassifications

Support Vector Machines (SVM)




No misclassifications—but is this the best position?

Support Vector Machines (SVM)




No misclassifications—but is this the best position?

Support Vector Machines (SVM)



Maximize the region between classes


Similarity Between Logistic Regression and SVM



Classification with SVMs

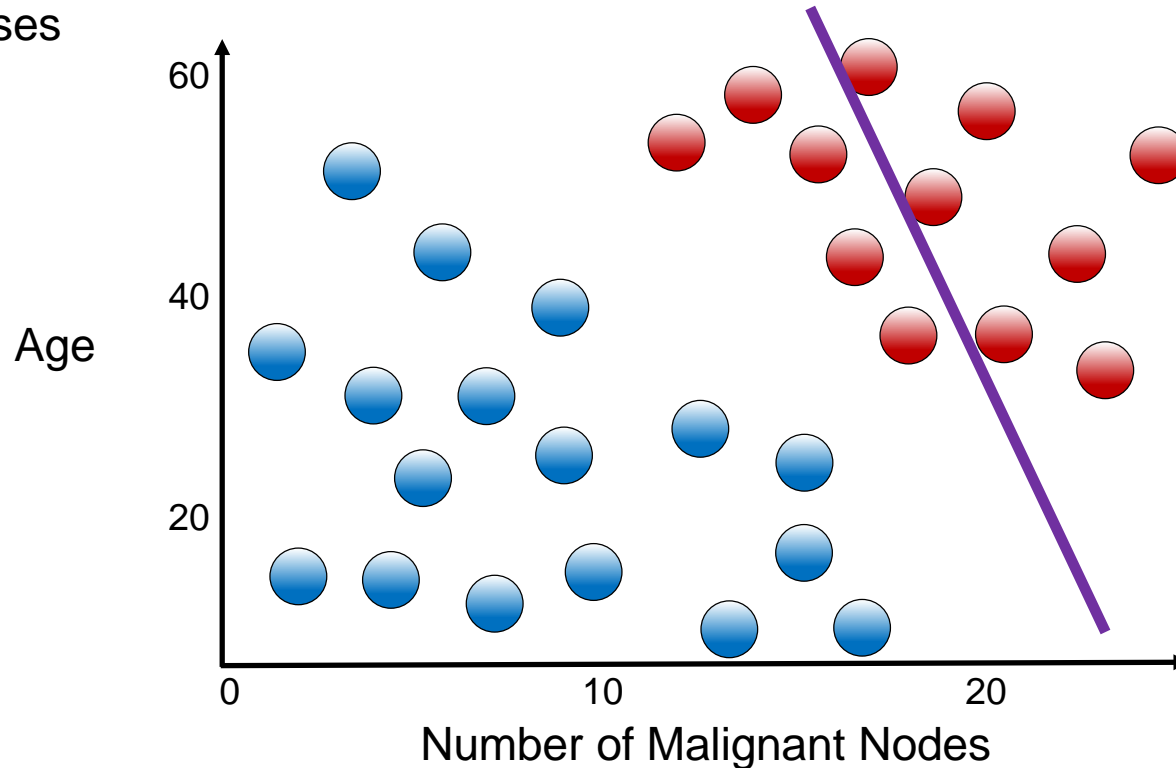
Two features (nodes, age)

Two labels (survived, lost)




Classification with SVMs

Find the line that best separates classes




Classification with SVMs

Find the line that best separates classes



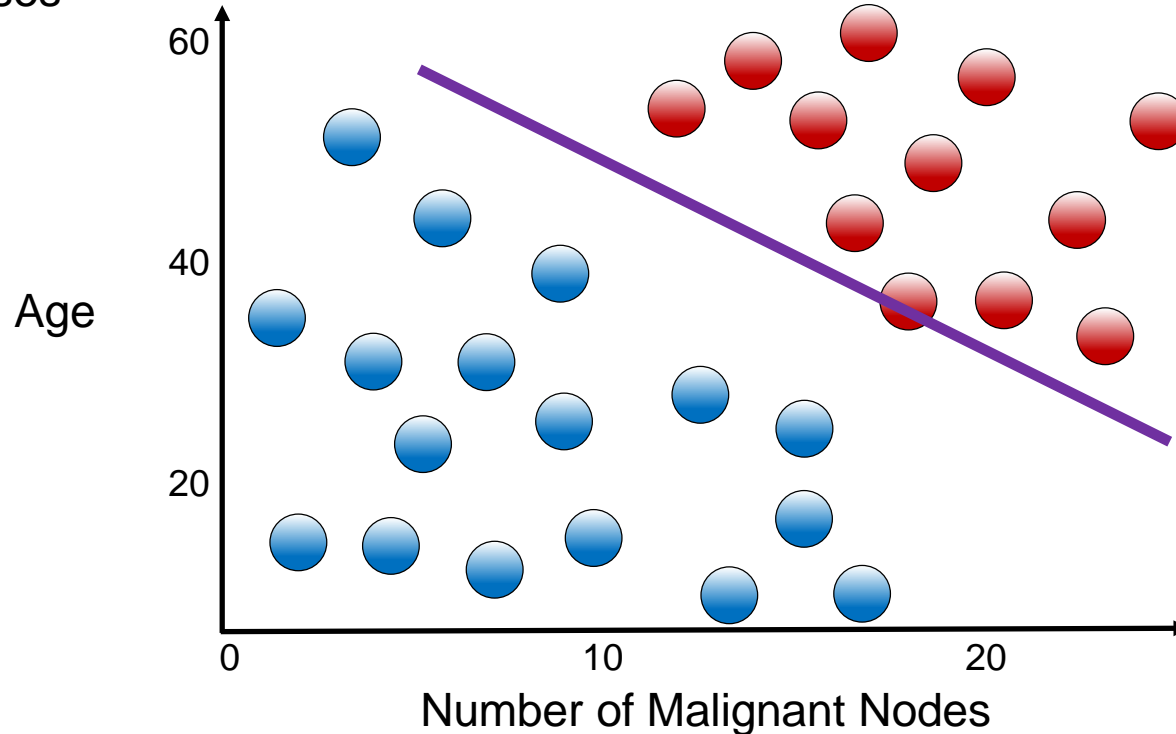
Classification with SVMs

Find the line that best separates classes



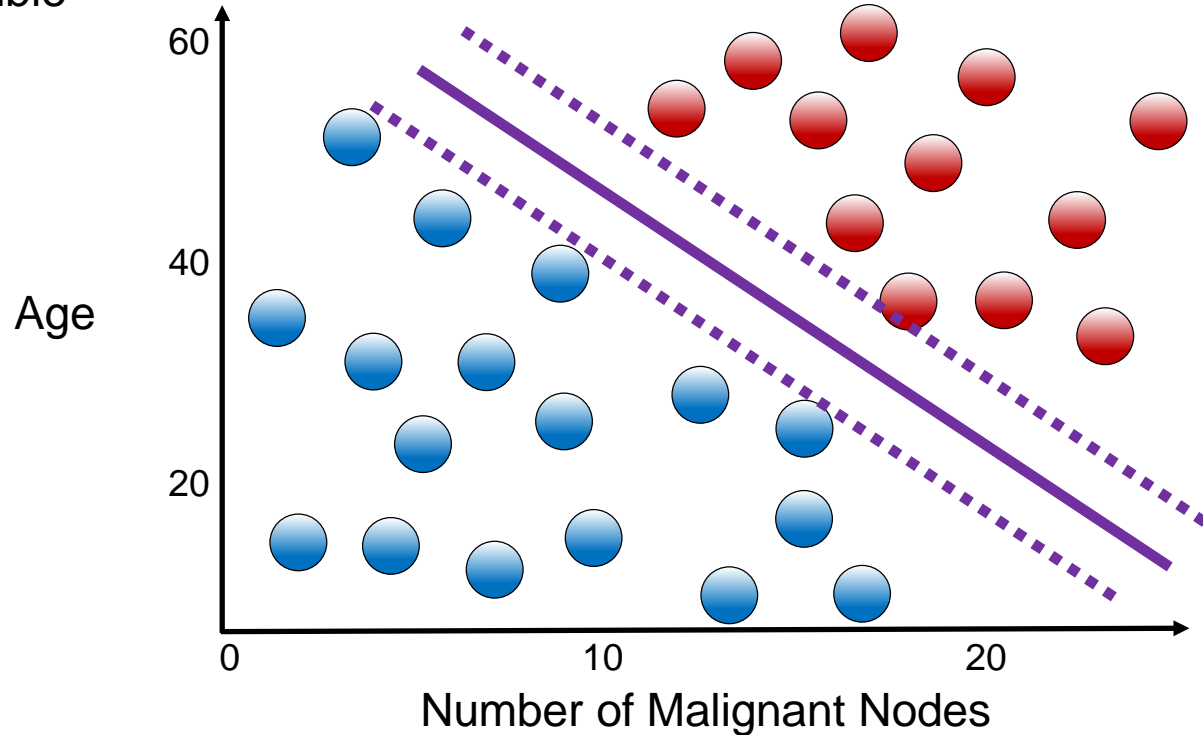
Classification with SVMs

Find the line that best separates classes




Classification with SVMs



And include the largest boundary possible






Logistic Regression vs SVM Cost Functions





Logistic Regression vs SVM Cost Functions





Logistic Regression vs SVM Cost Functions





The SVM Cost Function





The SVM Cost Function





The SVM Cost Function




The SVM Cost Function




The SVM Cost Function




Outlier Sensitivity in SVMs




Outlier Sensitivity in SVMs



Outlier Sensitivity in SVMs




Outlier Sensitivity in SVMs




Outlier Sensitivity in SVMs

This is probably still the correct boundary




Regularization in SVMs

$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$




Regularization in SVMs

$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$



Regularization in SVMs


$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$



Regularization in SVMs


$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$

← Large




Regularization in SVMs

$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$




Regularization in SVMs

$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$




Regularization in SVMs

$$J(\beta_i) = \text{SVMCost}(\beta_i) + \frac{1}{c} \sum_i \beta_i$$




Interpretation of SVM Coefficients

$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$



Interpretation of SVM Coefficients

$$J(\beta_i) = SVMCost(\beta_i) + \frac{1}{c} \sum_i \beta_i$$



Linear SVM: The Syntax

Import the class containing the classification method

```
from sklearn.svm import LinearSVC
```


Linear SVM: The Syntax

Import the class containing the classification method

```
from sklearn.svm import LinearSVC
```

Create an instance of the class

```
LinSVC = LinearSVC(penalty='l2', C=10.0)
```

Linear SVM: The Syntax

Import the class containing the classification method

```
from sklearn.svm import LinearSVC
```

Create an instance of the class

```
LinSVC = LinearSVC(penalty='l2', C=10.0)
```



regularization
parameters

Linear SVM: The Syntax

Import the class containing the classification method

```
from sklearn.svm import LinearSVC
```

Create an instance of the class

```
LinSVC = LinearSVC(penalty='l2', C=10.0)
```

Fit the instance on the data and then predict the expected value

```
LinSVC = LinSVC.fit(X_train, y_train)
```

```
y_predict = LinSVC.predict(X_test)
```

Linear SVM: The Syntax

Import the class containing the classification method

```
from sklearn.svm import LinearSVC
```

Create an instance of the class

```
LinSVC = LinearSVC(penalty='l2', C=10.0)
```

Fit the instance on the data and then predict the expected value


```
LinSVC = LinSVC.fit(X_train, y_train)
```

```
y_predict = LinSVC.predict(X_test)
```

Tune regularization parameters with cross-validation.


Kernels

Classification with SVMs




Non-Linear Decision Boundaries with SVM

Non-linear data can be made
linear with higher dimensionality




The Kernel Trick

Transform data so it is
linearly separable



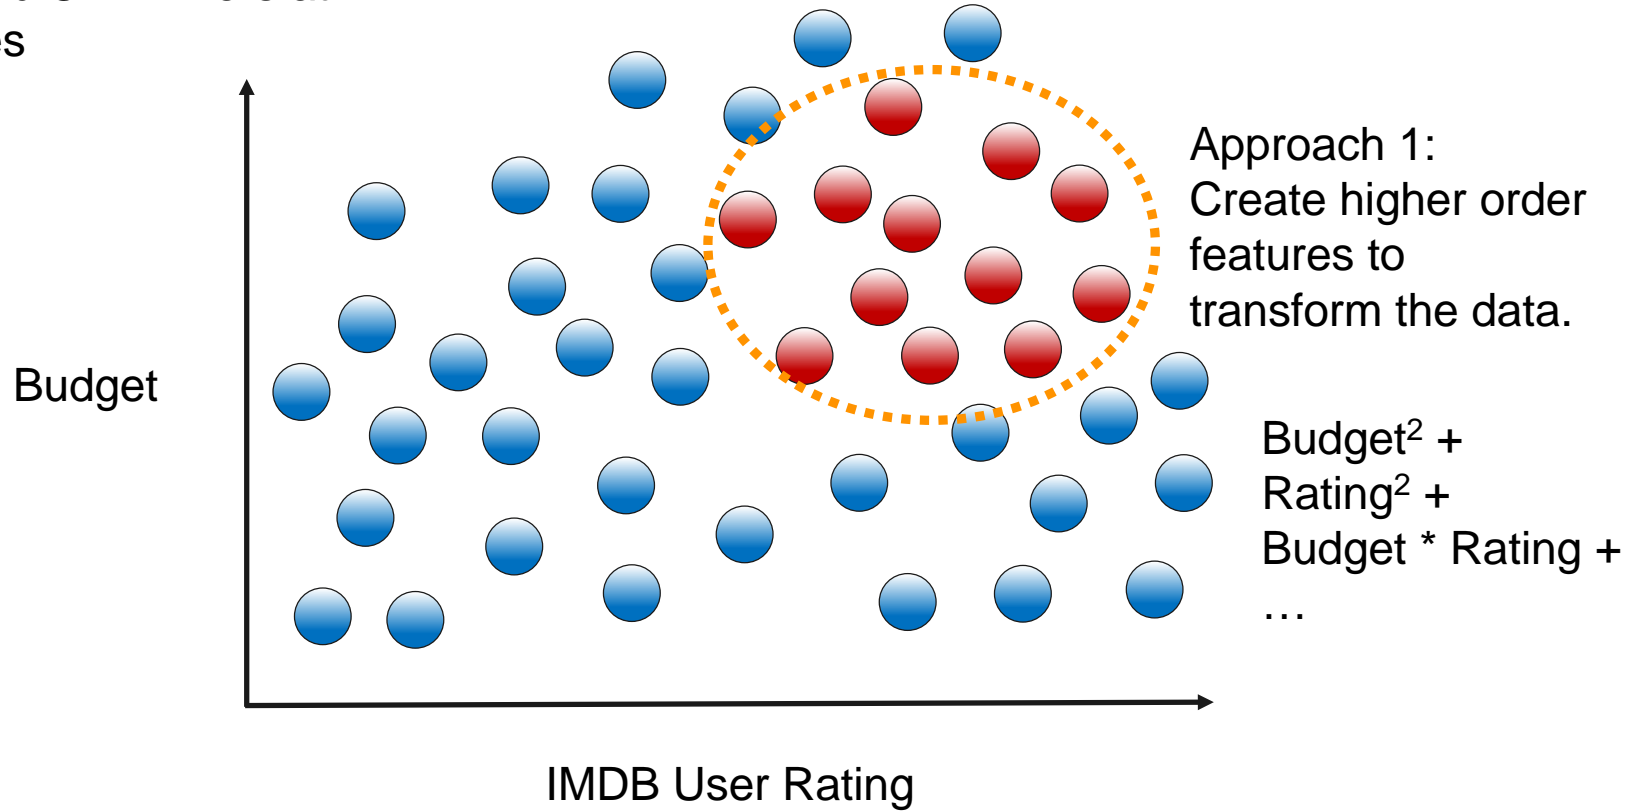
SVM Gaussian Kernel

Palme d'Or Winners at
Cannes




SVM Gaussian Kernel

Palme d'Or Winners at
Cannes




SVM Gaussian Kernel

Palme d'Or Winners at
Cannes




SVM Gaussian Kernel

Palme d'Or Winners at
Cannes




SVM Gaussian Kernel

Palme d'Or Winners at
Cannes



SVM Gaussian Kernel

Palme d'Or Winners at
Cannes




SVM Gaussian Kernel

Palme d'Or Winners at
Cannes

$$a_1(x^{obs}) = \exp \left[\frac{-\sum (x_i^{obs} - x_i^{Pulp Fiction})^2}{2\sigma^2} \right]$$

Budget



Create a
Gaussian function
at feature 1


IMDB User Rating

SVM Gaussian Kernel

Palme d'Or Winners at
Cannes

$$a_1(x^{obs}) = \exp \left[\frac{-\sum (x_i^{obs} - x_i^{Black\ Swan})^2}{2\sigma^2} \right]$$

Budget




Create a
Gaussian function
at feature 2

IMDB User Rating

SVM Gaussian Kernel

Palme d'Or Winners at
Cannes



Budget


IMDB User Rating

$$a_1(x^{obs}) = \exp \left[\frac{-\sum (x_i^{obs} - x_i^{Transformers})^2}{2\sigma^2} \right]$$




Create a
Gaussian function
at feature 3


SVM Gaussian Kernel



SVM Gaussian Kernel




SVM Gaussian Kernel



SVM Gaussian Kernel

Transformation:


$$[x_1, x_2] \rightarrow [0.7a_1, 0.9a_2, -0.6a_3]$$



Classification in the New Space


Transformation:

$$[x_1, x_2] \rightarrow [0.7\mathbf{a_1}, 0.9\mathbf{a_2}, -0.6\mathbf{a_3}]$$




SVM Gaussian Kernel

Palme d'Or Winners at
Cannes




SVM Gaussian Kernel

Palme d'Or Winners at
Cannes



SVM Gaussian Kernel

Palme d'Or Winners at
Cannes



SVMs with Kernels: The Syntax

Import the class containing the classification method

```
from sklearn.svm import SVC
```

SVMs with Kernels: The Syntax

Import the class containing the classification method

```
from sklearn.svm import SVC
```

Create an instance of the class

```
rbfSVC = SVC(kernel='rbf', gamma=1.0, C=10.0)
```

SVMs with Kernels: The Syntax

Import the class containing the classification method

```
from sklearn.svm import SVC
```

Create an instance of the class

```
rbfSVC = SVC(kernel='rbf', gamma=1.0, C=10.0)
```



set kernel and
associated
coefficient
(gamma)

»

SVMs with Kernels: The Syntax

Import the class containing the classification method

```
from sklearn.svm import SVC
```

Create an instance of the class

```
rbfSVC = SVC(kernel='rbf', gamma=1.0, C=10.0)
```



"C" is penalty
associated with
the error term

»

SVMs with Kernels: The Syntax

Import the class containing the classification method

```
from sklearn.svm import SVC
```

Create an instance of the class

```
rbfSVC = SVC(kernel='rbf', gamma=1.0, C=10.0)
```

Fit the instance on the data and then predict the expected value

```
rbfSVC = rbfSVC.fit(X_train, y_train)  
y_predict = rbfSVC.predict(X_test)
```

SVMs with Kernels: The Syntax

Import the class containing the classification method

```
from sklearn.svm import SVC
```

Create an instance of the class

```
rbfSVC = SVC(kernel='rbf', gamma=1.0, C=10.0)
```

Fit the instance on the data and then predict the expected value

```
rbfSVC = rbfSVC.fit(X_train, y_train)  
y_predict = rbfSVC.predict(X_test)
```

Tune kernel and associated parameters with cross-validation.

Feature Overload

Problem

SVMs with RBF Kernels are very slow to train with lots of features or data

Feature Overload

Problem

SVMs with RBF Kernels are very slow to train with lots of features or data

Solution

Construct approximate kernel map with SGD using Nystroem or RBF sampler

Feature Overload

Problem

SVMs with RBF Kernels are very slow to train with lots of features or data

Solution

Construct approximate kernel map with SGD using Nystroem or RBF sampler.
Fit a linear classifier.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method

```
from sklearn.kernel_approximation import Nystroem
```

Create an instance of the class

```
nystroemSVC = Nystroem(kernel='rbf', gamma=1.0,  
                        n_components=100)
```

Fit the instance on the data and transform

```
X_train = nystroemSVC.fit_transform(X_train)
```

```
X_test = nystroemSVC.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method

```
from sklearn.kernel_approximation import Nystroem
```

Create an instance of the class

```
nystroemSVC = Nystroem(kernel='rbf', gamma=1.0,  
                        n_components=100)
```



multiple non-linear kernels
can be used

Fit the instance on the data and transform

```
X_train = nystroemSVC.fit_transform(X_train)
```

```
X_test = nystroemSVC.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method

```
from sklearn.kernel_approximation import Nystroem
```

Create an instance of the class

```
nystroemSVC = Nystroem(kernel='rbf', gamma=1.0,  
                        n_components=100)
```



kernel and
gamma are
identical to SVC

Fit the instance on the data and transform

```
X_train = nystroemSVC.fit_transform(X_train)
```

```
X_test = nystroemSVC.transform(X_test)
```

Tune kernel parameters and components with cross-validation.


Faster Kernel Transformations: The Syntax

Import the class containing the classification method

```
from sklearn.kernel_approximation import Nystroem
```

Create an instance of the class

```
nystroemSVC = Nystroem(kernel='rbf', gamma=1.0,  
                        n_components=100)
```

 n_components
is number of
samples

Fit the instance on the data and transform

```
X_train = nystroemSVC.fit_transform(X_train)
```

```
X_test = nystroemSVC.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method

```
from sklearn.kernel_approximation import RBFsampler
```

Create an instance of the class

```
rbfSample = RBFsampler(gamma=1.0,  
                        n_components=100)
```

Fit the instance on the data and transform

```
X_train = rbfSample.fit_transform(X_train)  
X_test = rbfSample.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method

```
from sklearn.kernel_approximation import RBFsampler
```

Create an instance of the class

```
rbfSample = RBFsampler(gamma=1.0,  
                        n_components=100)
```



RBF is only
kernel that can
be used

Fit the instance on the data and transform

```
X_train = rbfSample.fit_transform(X_train)
```

```
X_test = rbfSample.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

Faster Kernel Transformations: The Syntax

Import the class containing the classification method

```
from sklearn.kernel_approximation import RBFsampler
```

Create an instance of the class

```
rbfSample = RBFsampler(gamma=1.0,  
                        n_components=100)
```



parameter
names are
identical to
previous

Fit the instance on the data and transform

```
X_train = rbfSample.fit_transform(X_train)
```

```
X_test = rbfSample.transform(X_test)
```

Tune kernel parameters and components with cross-validation.

When to Use Logistic Regression vs SVC

Features

Data

Model Choice

Many (~10K

Features)

Few (<100 Features)

Few (<100 Features)

Small (1K rows)

Medium (~10k rows)

Many (>100K Points)

Simple, Logistic or LinearSVC

SVC with RBF

Add features, Logistic, LinearSVC
or Kernel Approx.

Legal Notices and Disclaimers

This presentation is for informational purposes only. INTEL MAKES NO WARRANTIES, EXPRESS OR IMPLIED, IN THIS SUMMARY.

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. Check with your system manufacturer or retailer or learn more at [intel.com](https://www.intel.com).

This sample source code is released under the [Intel Sample Source Code License Agreement](#).

Intel and the Intel logo are trademarks of Intel Corporation in the U.S. and/or other countries.

*Other names and brands may be claimed as the property of others.

Copyright © 2017, Intel Corporation. All rights reserved.