


Chapter 12

Dimensionality Reduction


Curse of Dimensionality

- Theoretically, increasing features should improve performance




Curse of Dimensionality

- Theoretically, increasing features should improve performance
- In practice, too many features leads to worse performance



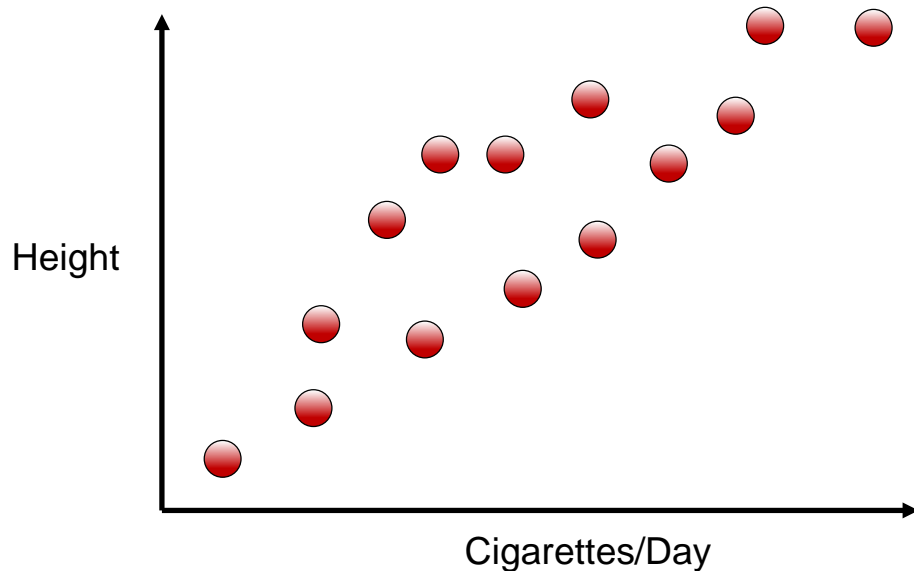
Curse of Dimensionality

- Theoretically, increasing features should improve performance
- In practice, too many features leads to worse performance
- Number of training examples required increases exponentially with dimensionality




Solution: Dimensionality Reduction

- Data can be represented by fewer dimensions (features)




Solution: Dimensionality Reduction

- Data can be represented by fewer dimensions (features)
- Reduce dimensionality by selecting subset (feature elimination)




Solution: Dimensionality Reduction

- Data can be represented by fewer dimensions (features)
- Reduce dimensionality by selecting subset (feature elimination)
- Combine with linear and non-linear transformations




Solution: Dimensionality Reduction

- Two features: height and cigarettes per day
- Both features increase together (correlated)
- Can we reduce number of features to one?




Solution: Dimensionality Reduction

- Two features: height and cigarettes per day
- Both features increase together (correlated)
- Can we reduce number of features to one?




Solution: Dimensionality Reduction

- Two features: height and cigarettes per day
- Both features increase together (correlated)
- Can we reduce number of features to one?




Solution: Dimensionality Reduction

- Two features: height and cigarettes per day
- Both features increase together (correlated)
- Can we reduce number of features to one?




Solution: Dimensionality Reduction

- Create single feature that is combination of height and cigarettes
- This is Principal Component Analysis (PCA)



Solution: Dimensionality Reduction

- Create single feature that is combination of height and cigarettes
- This is Principal Component Analysis (PCA)



Height + Cigarettes/Day


Dimensionality Reduction

Given an N -dimensional data set (x), find a $N \times K$ matrix (U):


$y = U^T x$, where y has K dimensions and $K < N$

$$x = \begin{bmatrix} x_1 \\ x_2 \\ \dots \\ x_n \end{bmatrix} \xrightarrow{U^T} y = \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_k \end{bmatrix} (K < N)$$


Principal Component Analysis (PCA)



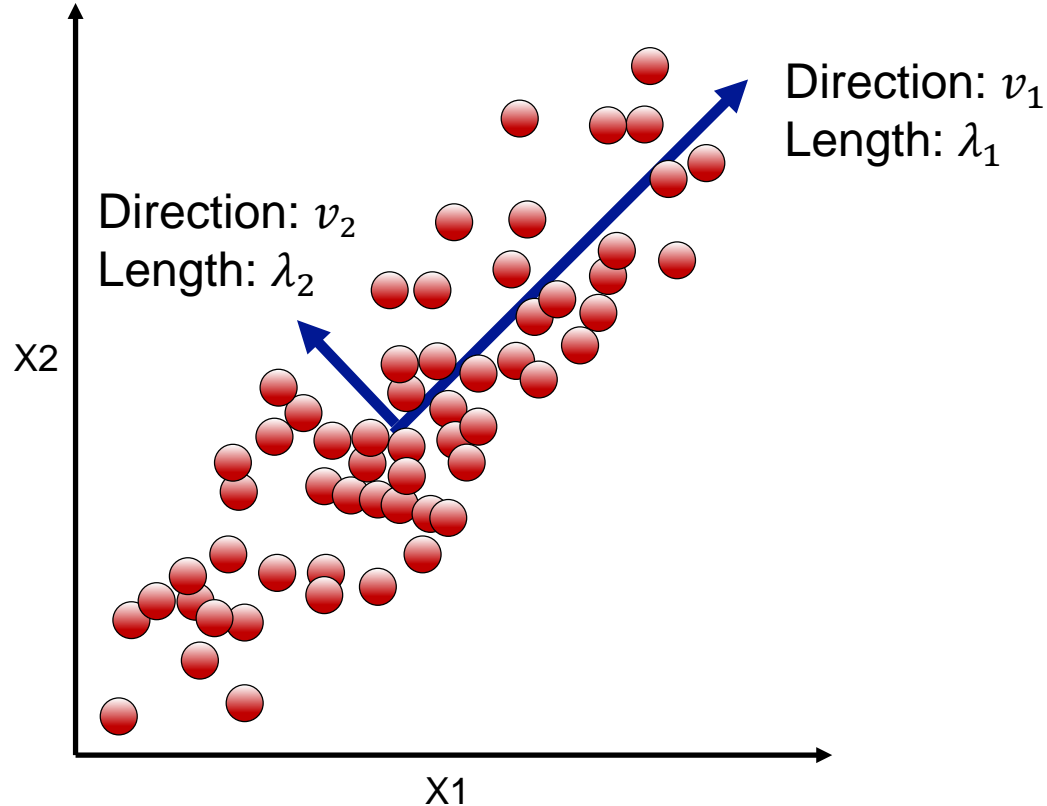
Principal Component Analysis (PCA)



Principal Component Analysis (PCA)



Principal Component Analysis (PCA)



Single Value Decomposition (SVD)

- SVD is a matrix factorization method normally used for PCA
- Does not require a square data set
- SVD is used by Scikit-learn for PCA

$$\begin{bmatrix} \star & \star & \star \\ \star & \star & \star \\ \star & \star & \star \\ \star & \star & \star \\ \star & \star & \star \end{bmatrix} = \begin{bmatrix} \star & \star & \star & \star & \star \\ \star & \star & \star & \star & \star \\ \star & \star & \star & \star & \star \\ \star & \star & \star & \star & \star \\ \star & \star & \star & \star & \star \end{bmatrix} \begin{bmatrix} \star & 0 & 0 \\ 0 & \star & 0 \\ 0 & 0 & \star \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} \star & \star & \star \\ \star & \star & \star \\ \star & \star & \star \end{bmatrix}$$

$A_{m \times n}$ $U_{m \times m}$ $S_{m \times n}$ $V_{n \times n}^T$

Truncated Single Value Decomposition


- How can SVD be used for dimensionality reduction?
- Principal components are calculated from US
- "Truncated SVD" used for dimensionality reduction ($n \rightarrow k$)

$$\begin{bmatrix} \star & \star & \star \\ \star & \star & \star \\ \star & \star & \star \\ \star & \star & \star \\ \star & \star & \star \end{bmatrix} \approx \begin{bmatrix} \star & \star & \star & \star & \star \\ \star & \star & \star & \star & \star \\ \star & \star & \star & \star & \star \\ \star & \star & \star & \star & \star \\ \star & \star & \star & \star & \star \end{bmatrix} \begin{bmatrix} 9 & 0 & 0 \\ 0 & 7 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} \star & \star & \star \\ \star & \star & \star \\ \star & \star & \star \\ \star & \star & \star \end{bmatrix}$$

$A_{m \times n}$
 $U_{m \times k}$
 $S_{k \times k}$
 $V_{k \times n}^T$

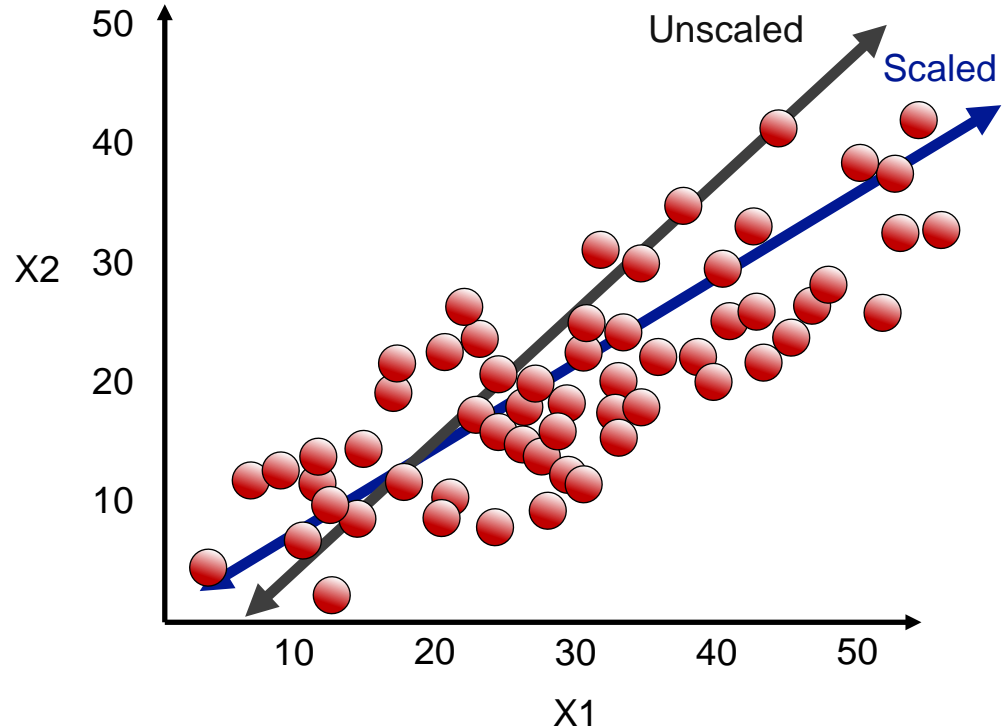
Importance of Feature Scaling

- PCA and SVD seek to find the vectors that capture the most variance
- Variance is sensitive to axis scale



Importance of Feature Scaling

- PCA and SVD seek to find the vectors that capture the most variance
- Variance is sensitive to axis scale
- Must scale data!



PCA: The Syntax

Import the class containing the dimensionality reduction method

`from sklearn.decomposition import PCA`

PCA: The Syntax

Import the class containing the dimensionality reduction method

```
from sklearn.decomposition import PCA
```

Create an instance of the class

```
PCAinst = PCA(n_components=3, whiten=True)
```


PCA: The Syntax

Import the class containing the dimensionality reduction method

```
from sklearn.decomposition import PCA
```

Create an instance of the class

```
PCAinst = PCA(n_components=3, whiten=True)
```



final number of
dimensions

PCA: The Syntax

Import the class containing the dimensionality reduction method

```
from sklearn.decomposition import PCA
```

Create an instance of the class

```
PCAinst = PCA(n_components=3, whiten=True)
```



whiten = scale
and center data

PCA: The Syntax

Import the class containing the dimensionality reduction method

```
from sklearn.decomposition import PCA
```

Create an instance of the class

```
PCAinst = PCA(n_components=3, whiten=True)
```

Fit the instance on the data and then transform the data

```
X_trans = PCAinst.fit_transform(X_train)
```

PCA: The Syntax

Import the class containing the dimensionality reduction method

```
from sklearn.decomposition import PCA
```

Create an instance of the class

```
PCAinst = PCA(n_components=3, whiten=True)
```

Fit the instance on the data and then transform the data

```
X_trans = PCAinst.fit_transform(X_train)
```

Does not work with sparse matrices

Truncated SVD: The Syntax

Import the class containing the dimensionality reduction method

```
from sklearn.decomposition import TruncatedSVD
```

Create an instance of the class

```
SVD = TruncatedSVD(n_components=3)
```

Fit the instance on the data and then transform the data

```
X_trans = SVD.fit_transform(X_sparse)
```

Works with sparse matrices—used with text data for Latent Semantic Analysis (LSA)

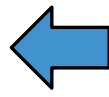
Truncated SVD: The Syntax

Import the class containing the dimensionality reduction method

```
from sklearn.decomposition import TruncatedSVD
```

Create an instance of the class

```
SVD = TruncatedSVD(n_components=3)
```

 does not
center data


Fit the instance on the data and then transform the data

```
X_trans = SVD.fit_transform(X_sparse)
```

Works with sparse matrices—used with text data for Latent Semantic Analysis (LSA)

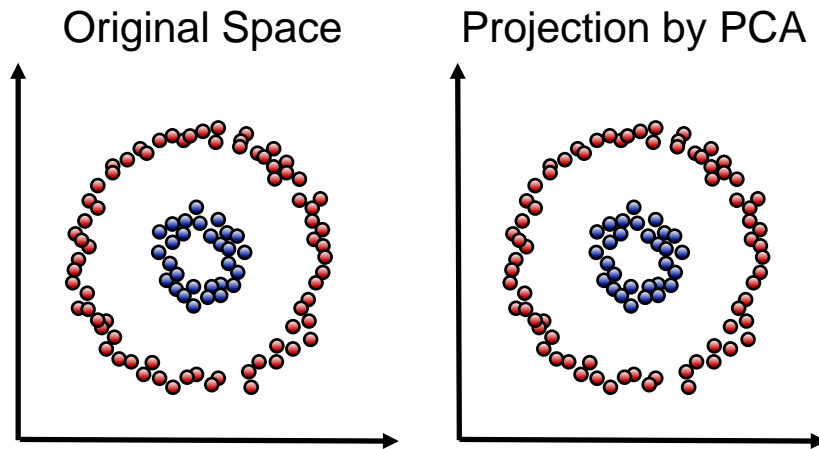
Moving Beyond Linearity

- Transformations calculated with PCA/SVD are linear
- Data can have non-linear features




Moving Beyond Linearity

- Transformations calculated with PCA/SVD are linear
- Data can have non-linear features
- This can cause dimensionality reduction to fail



Moving Beyond Linearity


- Transformations calculated with PCA/SVD are linear
- Data can have non-linear features
- This can cause dimensionality reduction to fail




Kernel PCA

- **Solution:** kernels can be used to perform non-linear PCA

Original Space





Projection by KPCA



Kernel PCA

- **Solution:** kernels can be used to perform non-linear PCA
- Like the kernel trick introduced for SVMs



Kernel PCA: The Syntax

Import the class containing the dimensionality reduction method

```
from sklearn.decomposition import KernelPCA
```

Create an instance of the class


```
kPCA = KernelPCA(n_components=3, kernel='rbf', gamma=1.0)
```

Fit the instance on the data and then transform the data

```
X_trans = kPCA.fit_transform(X_train)
```

Multi-Dimensional Scaling (MDS)

- Non-linear transformation
- Doesn't focus on maintaining overall variance
- Instead, maintains geometric distances between points



MDS: The Syntax

Import the class containing the dimensionality reduction method

```
from sklearn.manifold import MDS
```

Create an instance of the class

```
mdsMod = MDS(n_components=2)
```

Fit the instance on the data and then transform the data

```
X_trans = mdsMod.fit_transform(X_sparse)
```

Many other manifold dimensionality methods exist: **Isomap**, **TSNE**.

Uses of Dimensionality Reduction

- Frequently used for high dimensionality data
- Natural language processing (NLP)—many word combinations
- Image-based data sets—pixels are features



Image Source: https://commons.wikimedia.org/wiki/File:Monarch_In_May.jpg

Uses of Dimensionality Reduction

- Divide image into 12 x 12 pixel sections




Image Source: https://commons.wikimedia.org/wiki/File:Monarch_In_May.jpg

Uses of Dimensionality Reduction

- Divide image into 12 x 12 pixel sections
- Flatten section to create row of data with 144 features

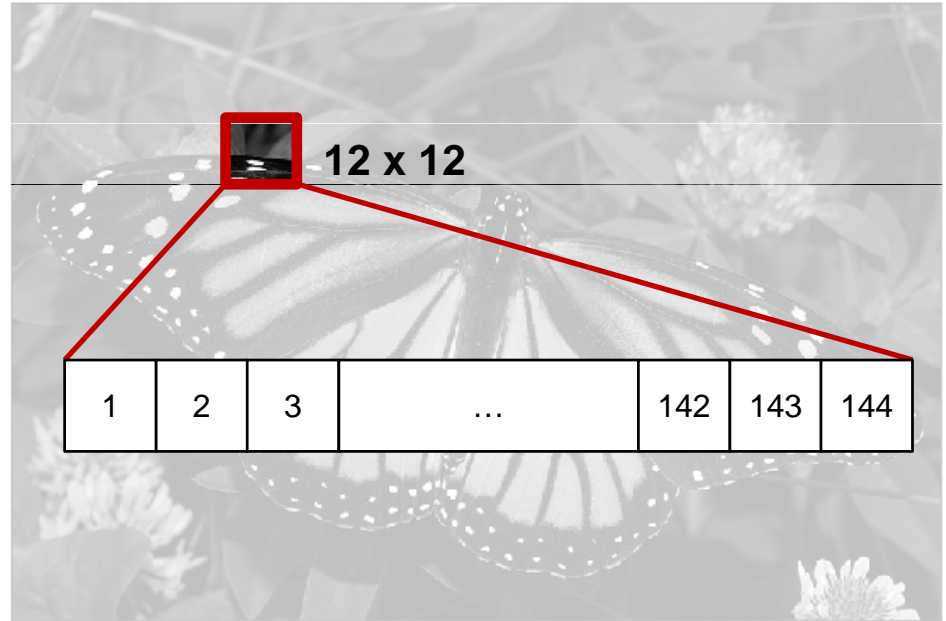
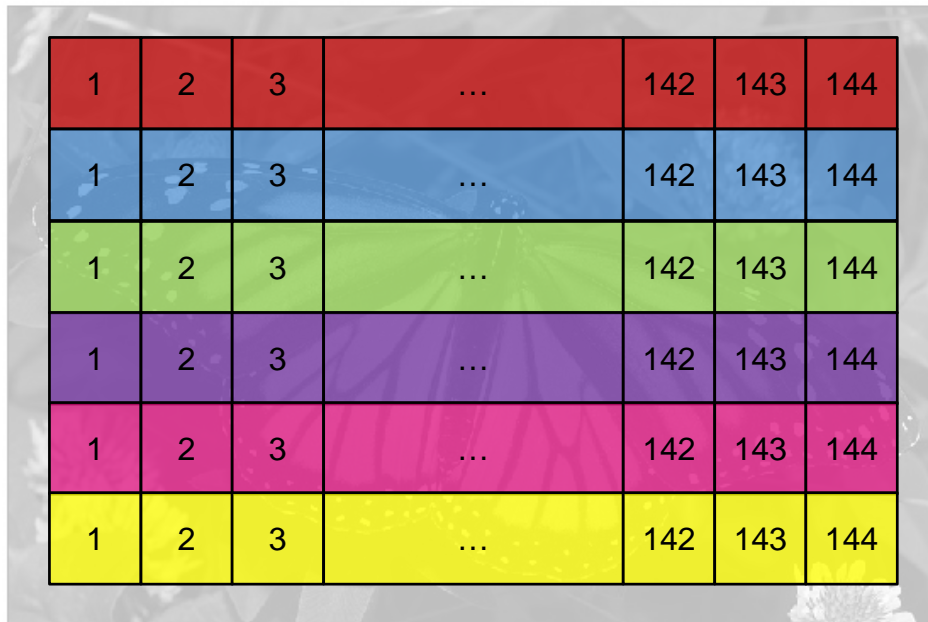


Image Source: https://commons.wikimedia.org/wiki/File:Monarch_In_May.jpg

Uses of Dimensionality Reduction

- Divide image into 12 x 12 pixel sections
- Flatten section to create row of data with 144 features
- Perform PCA on all data points



1	2	3	...	142	143	144
1	2	3	...	142	143	144
1	2	3	...	142	143	144
1	2	3	...	142	143	144
1	2	3	...	142	143	144
1	2	3	...	142	143	144

Image Source: https://commons.wikimedia.org/wiki/File:Monarch_In_May.jpg

PCA Compression: 144 \rightarrow 60 Dimensions



Dimensions



Dimensions

PCA Compression: 144 \rightarrow 16 Dimensions




Dimensions



16 Dimensions

Sixteen Most Important Eigenvectors



PCA Compression: 144 \rightarrow 4 Dimensions




144

Dimensions




4 Dimensions


L2 Error and PCA Dimension



Four Most Important Eigenvectors



Four Most Important Eigenvectors




PCA Compression: $144 \rightarrow 1$ Dimension



144

Dimensions



1

Dimension

Legal Notices and Disclaimers

This presentation is for informational purposes only. INTEL MAKES NO WARRANTIES, EXPRESS OR IMPLIED, IN THIS SUMMARY.

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. Check with your system manufacturer or retailer or learn more at [intel.com](https://www.intel.com).

This sample source code is released under the [Intel Sample Source Code License Agreement](#).

Intel and the Intel logo are trademarks of Intel Corporation in the U.S. and/or other countries.

*Other names and brands may be claimed as the property of others.

Copyright © 2017, Intel Corporation. All rights reserved.